



Bastian Wittmann PhD candidate

Research Interest: Graph Neural Networks, Vision Transformers, Blood Vessel Graphs
Google Scholar: <https://scholar.google.com/citations?user=0mxtU-cAAAAJ&hl=en&oi=ao>
GitHub: <https://github.com/bwittmann>
LinkedIn: <https://ch.linkedin.com/in/bastian-wittmann-7830731b4>

Medical Image **Classification** and Segmentation with Deep Learning - Part 1

Prof. Dr. Bjoern Menze

Bastian Wittmann

Biomedical Image Analysis and Machine Learning @ UZH

Agenda

PART 1: Medical Image Classification - Bastian Wittmann

- 1) Brief Introduction to Machine Learning and Deep Learning
- 2) E1: Hands-On Session
 - a) Introduction to PyTorch (see notebook section 0)
 - b) Preparing Data (see notebook section 1)
 - c) Implementation of a Fully Connected Network (see notebook section 2)
- 3) Convolutional Neural Networks and ResNet
- 4) E2: Hands-On Session
 - a) Implementation of a simple Convolutional Neural Network (see notebook section 3)
 - b) Experiments with Deep Residual Networks (see notebook section 4)
 - c) Optional: Interpretability of Predictions (see notebook section 5)

12am - 1pm: Lunch Break

PART 2: Medical Image Segmentation - Paul Bueschl

4pm: End of Practical Session

Colab Notebooks

- E1 [0) Introduction to PyTorch
- 1) Preparing the Data
- 2) Model 1 - Fully Connected Network (FCN/MLP)
- E2 [3) Model 2 - Simple Convolutional Neural Network (CNN)
- 4) Model 3 - Deep Residual Neural Network (ResNet)
- 5) Interpretability of Predictions

****Please ask whenever you have any open questions!****

We provide **two** versions of our notebook

- 1) version without solutions: [here](#) (recommended version)
- 2) version with solutions: [here](#) (solely if you feel super overwhelmed*)

**This should be considered the 'last resort'. Please ask us for help first.*

Big Picture

various different
input formats



...


Inputs

blackbox


Neural Network

Outputs

various different
output formats



dog: 0.9



...

Classification

x-ray image of lung



x_i

Inputs

blackbox θ

Neural Network

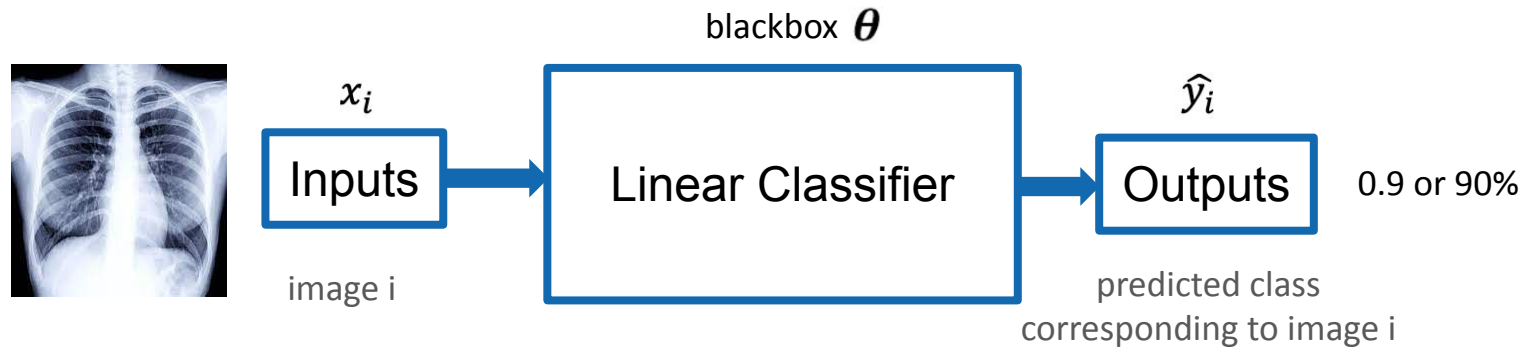
\hat{y}_i

Outputs

predict whether
patient has bacterial
pneumonia

0.9 or 90%

Linear Classification



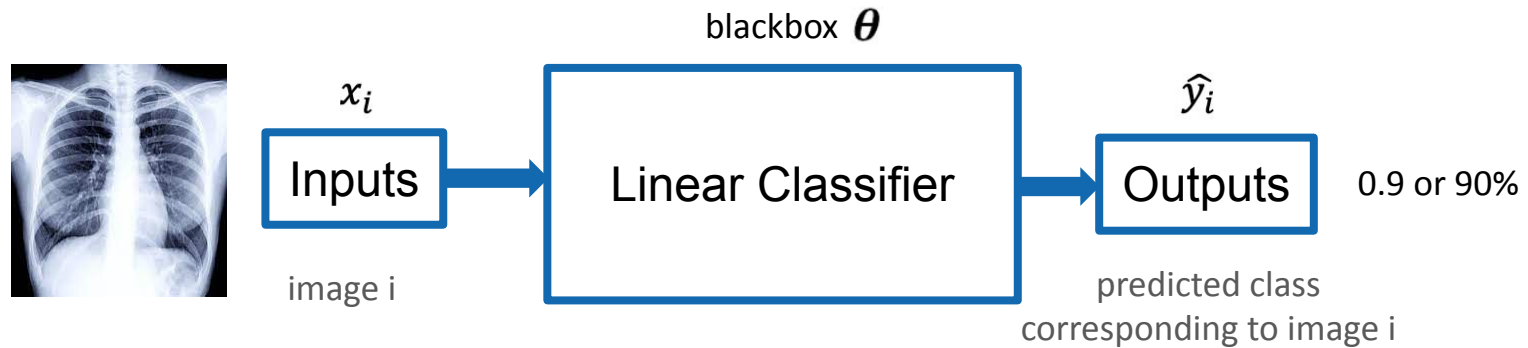
how to divide image
into different features
that we can use for
classification?

how can we create a
model representative of
our neural network?

how can we ensure
that our network
delivers predictions in
the range of $[0, 1]$?

three main questions we have to solve
for any classifier


Linear Classification




how to divide image
into different features
that we can use for
classification?

how can we create a
model representative of
our neural network?


how can we ensure
that our network
delivers predictions in
the range of $[0, 1]$?




Linear Classification



how to divide image
into different features
that we can use for
classification?

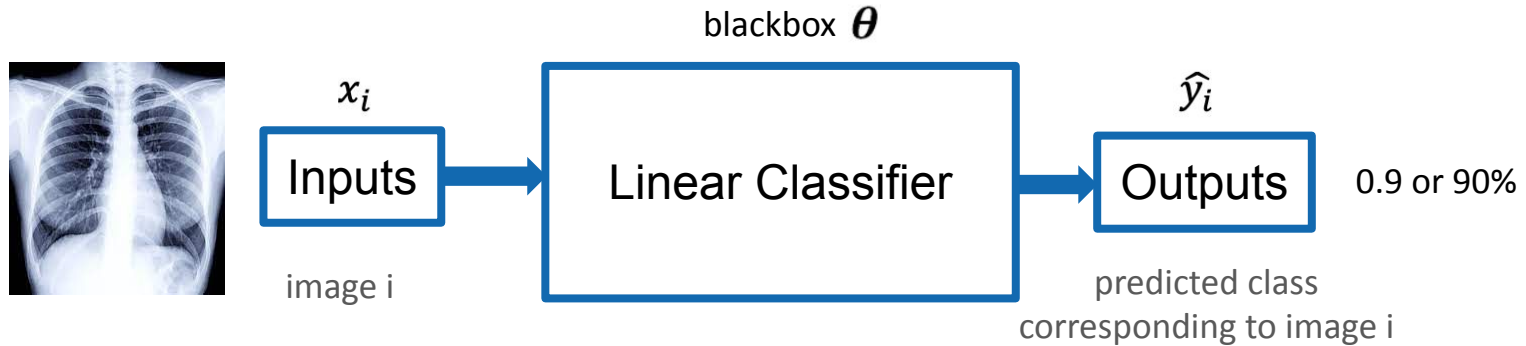


how can we create a
model representative of
our neural network?




how can we ensure
that our network
delivers predictions in
the range of $[0, 1]$?


Linear Classification



how to divide image
into different features
that we can use for
classification?




how can we create a
model representative of
our neural network?




*single layer
perceptron*

how can we ensure
that our network
delivers predictions in
the range of $[0, 1]$?


Linear Classification



how to divide image
into different features
that we can use for
classification?




how can we create a
model representative of
our neural network?

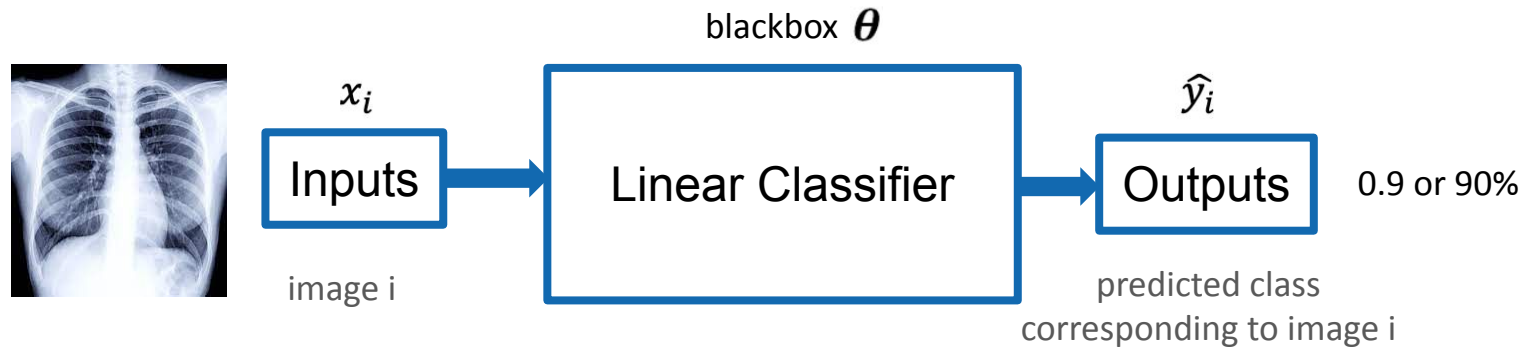


single layer
perceptron

how can we ensure
that our network
delivers predictions in
the range of $[0, 1]$?




Linear Classification



how to divide image
into different features
that we can use for
classification?

how can we create a
model representative of
our neural network?


how can we ensure
that our network
delivers predictions in
the range of $[0, 1]$?



Lets now put all of this
together to create our first
classifier


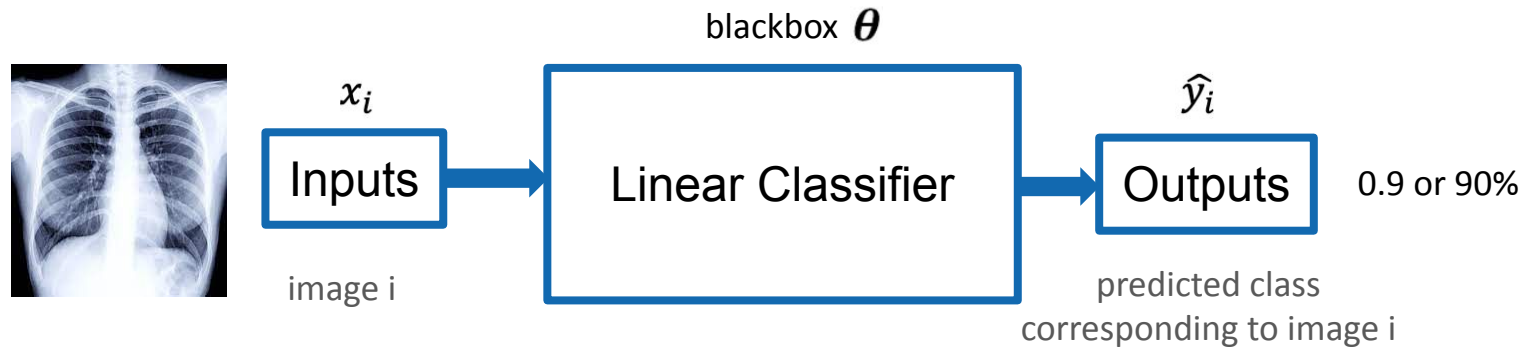
x_{id}

single layer
perceptron





$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

Linear Classification



*Logistic Regression:
a 1 layer, 1 neuron neural network*

Linear Classification




Logistic Regression:
a 1 layer, 1 neuron neural network

$$\hat{y}_i = \sigma(x_i \theta + \theta_0)$$


$$\theta_0 + \sum_{j=1}^d x_{ij} \theta_j = \theta_0 + x_{i1} \theta_1 + x_{i2} \theta_2 + \dots + x_{id} \theta_d$$

what are now all of
these parameters Θ ?

Linear Classification



the models parameters Θ are what we can actually modify to obtain better performances for our classification task!



Logistic Regression:
a 1 layer, 1 neuron neural network

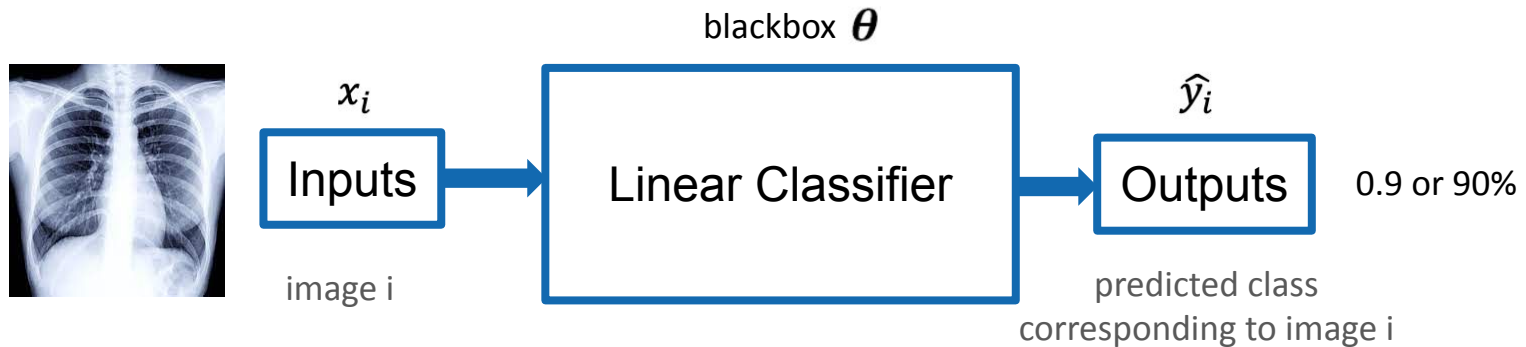
$$\hat{y}_i = \sigma(x_i \theta + \theta_0)$$

$$\theta_0 + \sum_{j=1}^d x_{ij} \theta_j = \theta_0 + x_{i1} \theta_1 + x_{i2} \theta_2 + \dots + x_{id} \theta_d$$


this is what is commonly referred to as **training** a neural network.

more on this later...

Linear Classification



this is, therefore,
typically referred to
as a **linear layer**




Logistic Regression:
a 1 layer, 1 neuron neural network

$$\hat{y}_i = \sigma(x_i \theta + \theta_0)$$


$$\theta_0 + \sum_{j=1}^d x_{ij} \theta_j = \theta_0 + x_{i1} \theta_1 + x_{i2} \theta_2 + \dots + x_{id} \theta_d$$

Linear Classification



*Logistic Regression:
a 1 neuron neural network*

Linear Classification

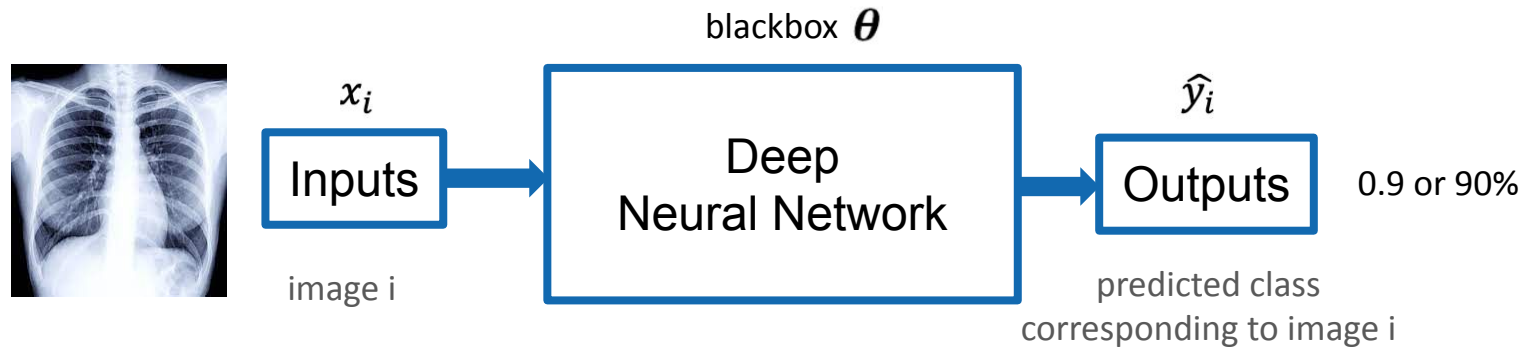


A linear classifier can solely learn a linear decision boundary!


Logistic Regression:

a 1 neuron neural network

Neural Networks



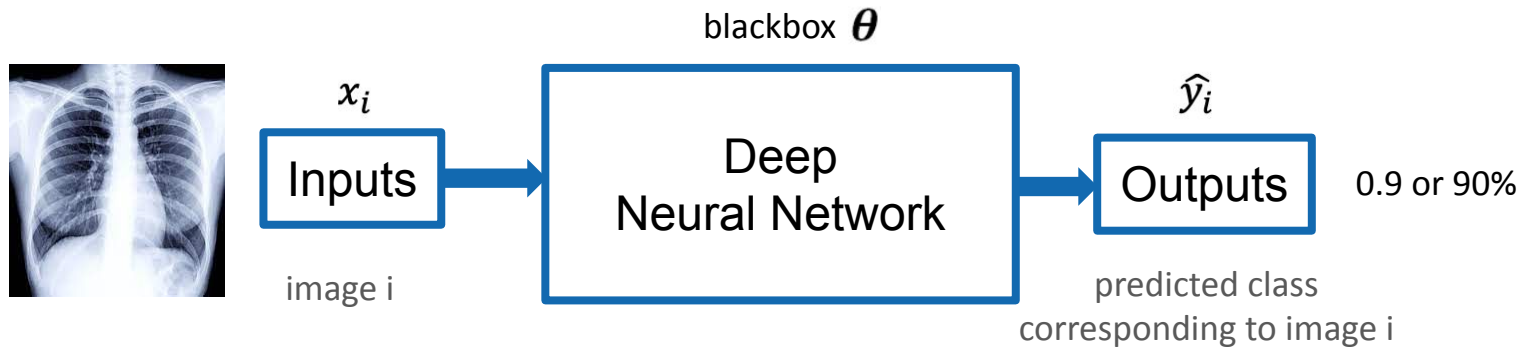
$$\hat{y}_i = \sigma(x_i \theta)$$




how can we
transform our simple
linear classifier to
learn complex
non-linear functions?

*Logistic Regression:
a 1 layer, 1 neuron neural network*

Neural Networks



$$\hat{y}_i = \sigma(x_i \theta)$$



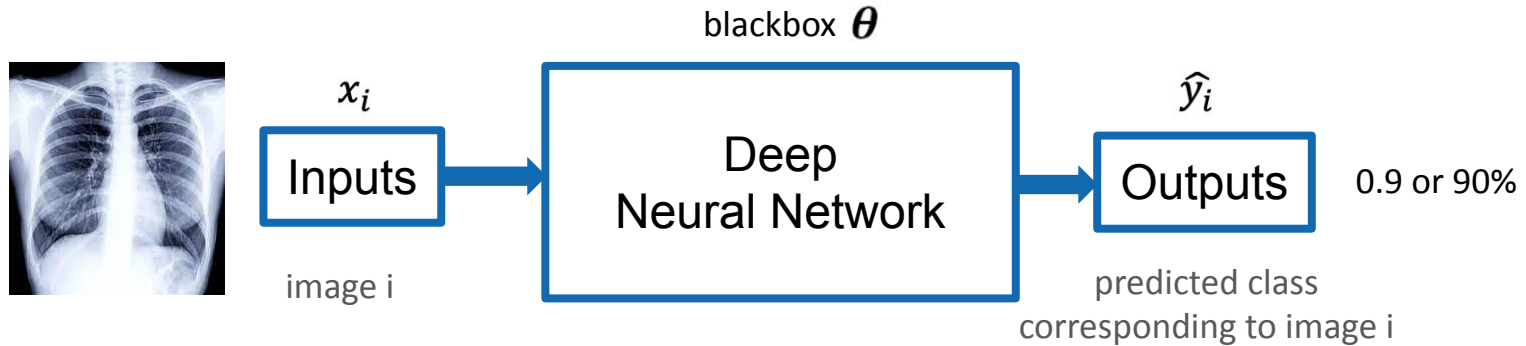
Logistic Regression:
a 1 layer, 1 neuron neural network

separate the neural network's
parameters via non-linearities to
enable us to learn a more
complex classifier!

$$\hat{y}_i = \sigma(x_i \theta) \longrightarrow \hat{y}_i = \sigma(\sigma(x_i \theta) \theta)$$

Non-linear Classifier:
a 2 layer neural network

Neural Networks




$$\hat{y}_i = W_5 \sigma(W_4 \tanh(W_3, \max(0, W_2 \max(0, W_1 x_i))))$$


lernable
parameters θ

a neural network is comprised of
multiple layers separated by
non-linearities

this results in **deep neural networks**
(that's why it's called deep learning)



Neural Networks




there exists a whole zoo of
non-linearities (or activations).

ReLU is typically the standard choice.
(best convergence, no vanishing gradients)


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$



ReLU


$$\max(0, x)$$



rectified linear unit

Leaky ReLU

$$\max(0.1x, x)$$



Maxout


$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



The Dataset

 $\{x_{1:n}, y_{1:n}\}$

whole annotated dataset (images + ground truth)

first, we split the
dataset into three
splits

the training,
validation, and
test set

60%

train

20%

validation

20%

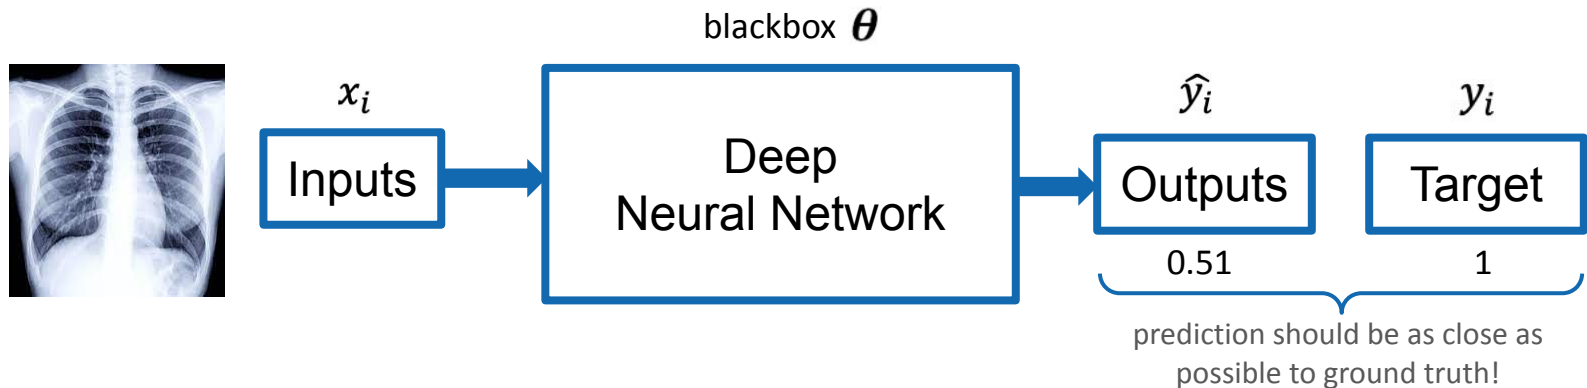
test

train the neural network's parameters θ

validate our design choices test performance on hidden test set

we don't use this data to optimize the parameters!!!

How to Train the Neural Network



we need a way to mathematically describe how close our predictions are to the target!

all images in our train data

binary cross-entropy (BCE):

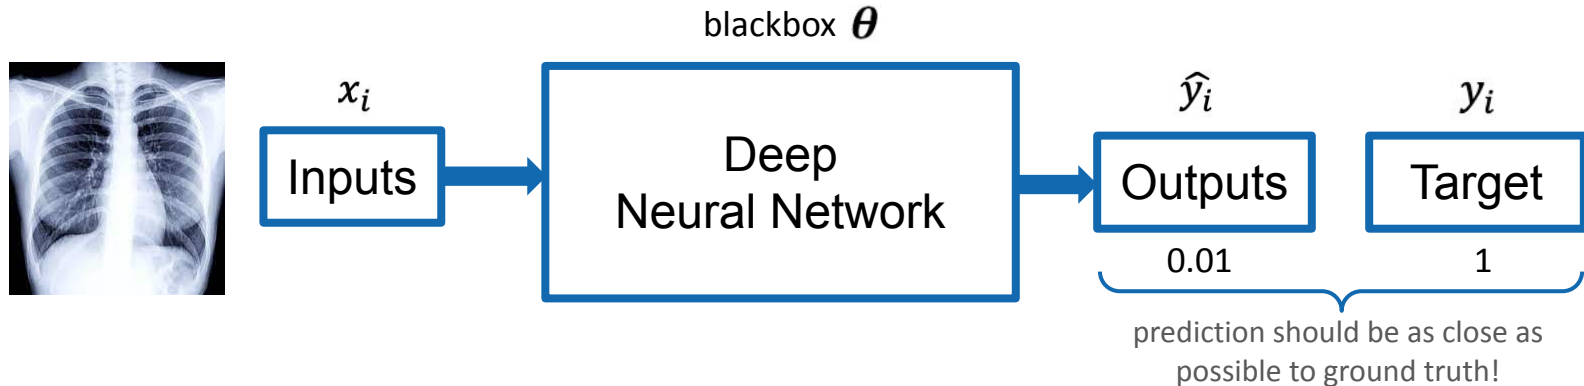
$$L(\mathbf{y}, \hat{\mathbf{y}}; \boldsymbol{\theta}) = - \sum_{i=1}^n (y_i \cdot \log \hat{y}_i + (1 - y_i) \cdot \log[1 - \hat{y}_i])$$

train loss

prediction of image i

ground truth

How to Train the Neural Network




we want a **large loss** for **bad performances**

here we predict the complete opposite

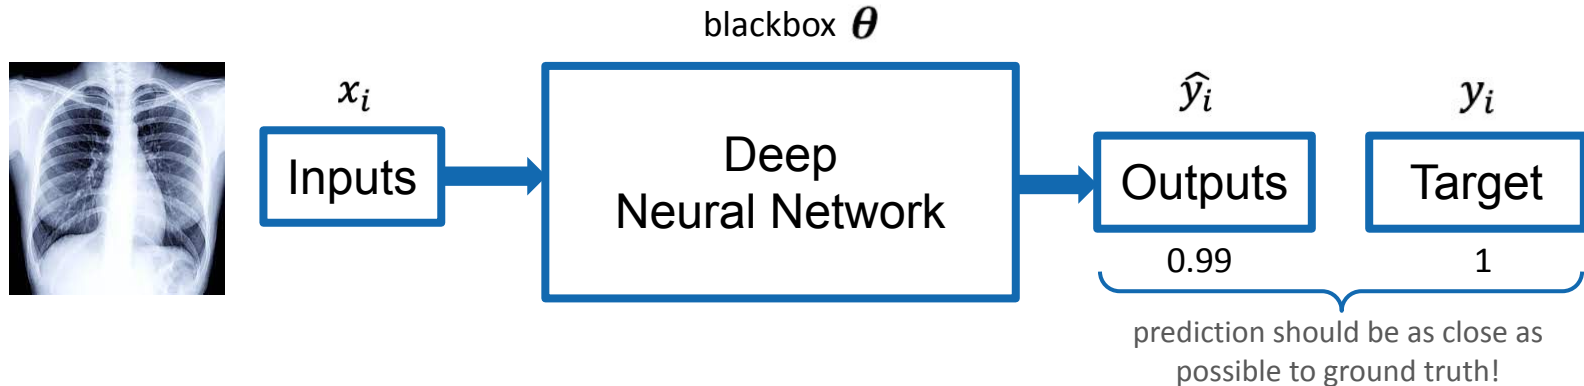
this should be penalized with a high loss

$$\mathcal{L}(\hat{y}_i, y_i) = \overbrace{-[y_i \log \hat{y}_i] + (1 - y_i) \log(1 - \hat{y}_i)}^{2.0000}$$

\uparrow \uparrow
 1 0.01



How to Train the Neural Network




we want a **small loss** for **good performances**

here we predict correct


this should be rewarded with a low loss

$$\mathcal{L}(\hat{y}_i, y_i) = \overbrace{-[y_i \log \hat{y}_i] + (1 - y_i) \log(1 - \hat{y}_i)}^{0.0044}$$

\uparrow \uparrow
 1 0.99



How to Train the Neural Network



the smaller the loss, the better the performance of our classifier on the train set

$$L(\mathbf{y}, \hat{\mathbf{y}}; \theta) = - \sum_{i=1}^n (y_i \cdot \log \hat{y}_i + (1 - y_i) \cdot \log[1 - \hat{y}_i])$$


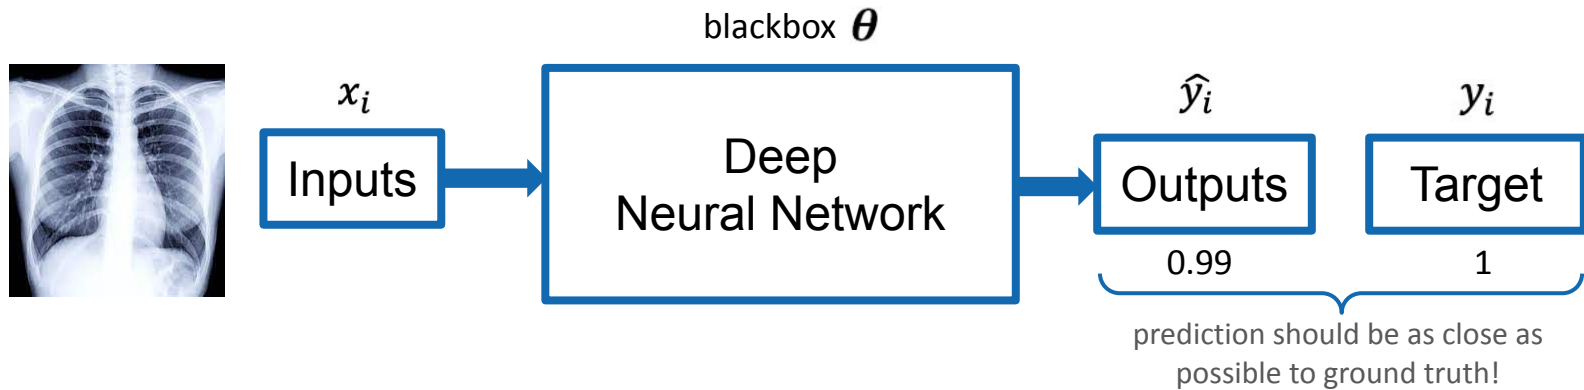
⇒ we aim to minimize the loss with respect to the trainable parameters Θ

$$\theta^* = \arg \min_{\theta} L(\mathbf{y}, \hat{\mathbf{y}}; \theta)$$


we **train our NN** by updating its parameters to fit the task at hand

However, **no closed-form solution** ⇒ iterative optimization methods: stochastic gradient descent


How to Train the Neural Network



How to Train the Neural Network



we iteratively minimize the loss $L(y, f_{\theta}(x))$ wrt. θ using **stochastic gradient descent**!



learning rate (each time step we just make a small update)



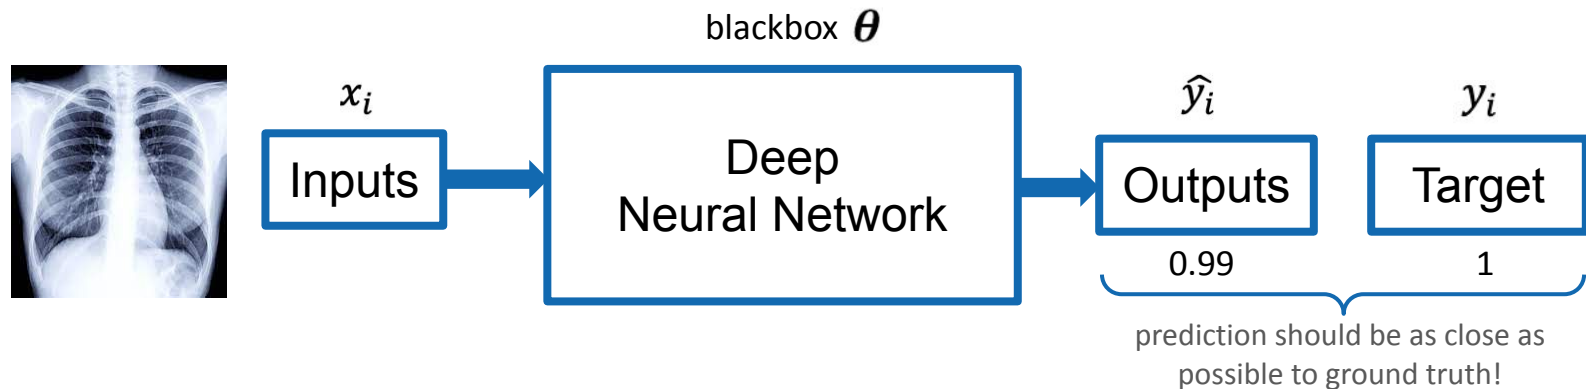
$$\theta_{t+1} = \theta_t - \alpha \nabla_{\theta} L(y, f_{\theta_t}(x))$$

we iteratively update our parameters until convergence

$$\nabla_{\theta} L = \frac{1}{m} \sum_{i=1}^m \nabla_{\theta} L_i$$


we approximate the gradient over a mini-batch for faster convergence (stochastic gradient descent)

How to Train the Neural Network




$$\theta^* = \arg \min_{\theta} L(y, \hat{y}; \theta)$$

How to Train the Neural Network

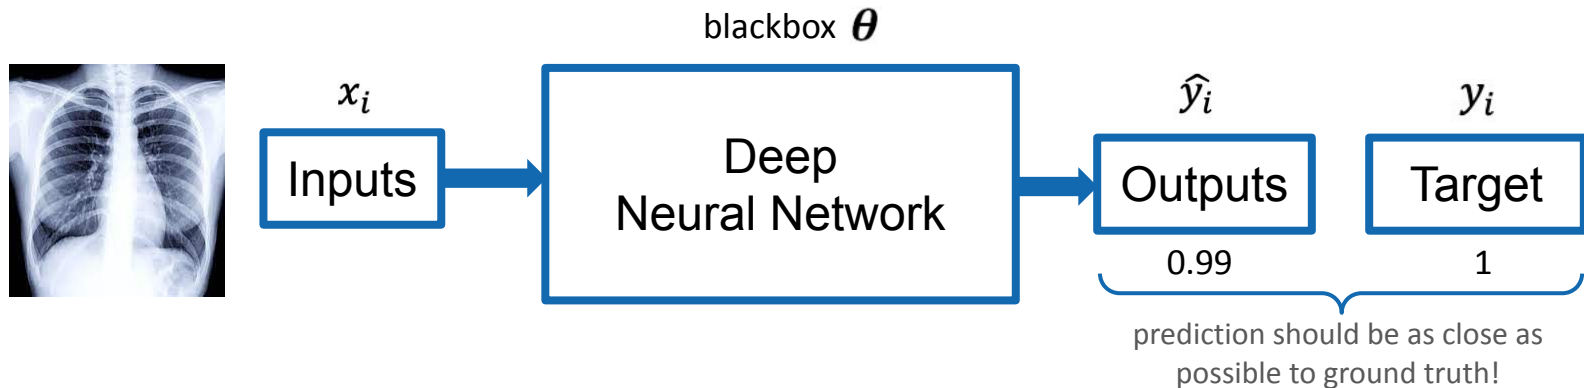


please remember that we split our dataset:




to ensure that our model does not just **memorize samples** seen during training, we always have to check whether the **model generalizes to data from the validation set**.

How to Train the Neural Network



to **avoid overfitting (reduce the generalization gap)** of our model to the training dataset we typically utilize **regularization** techniques:

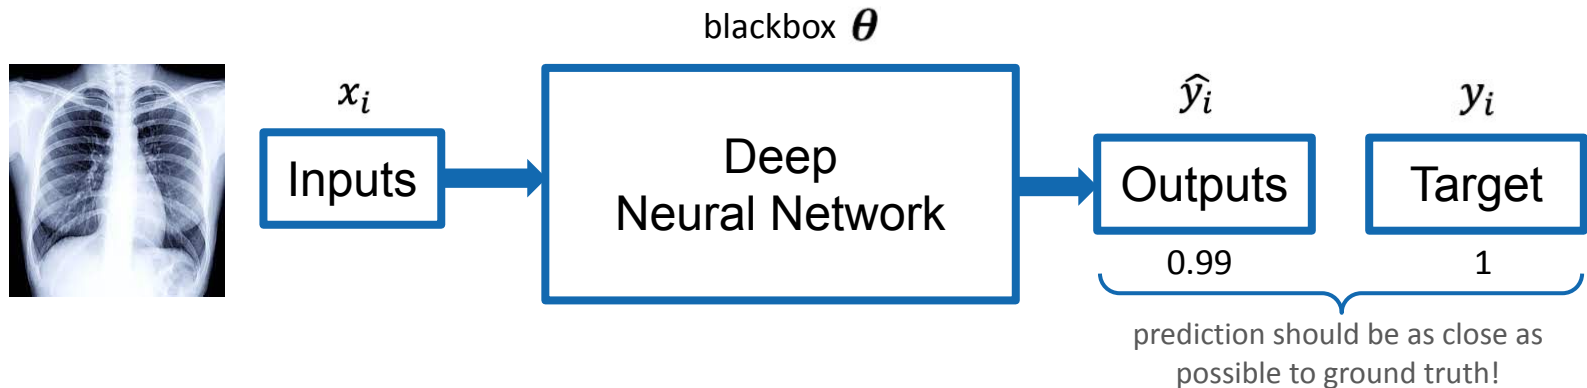


lower
validation error

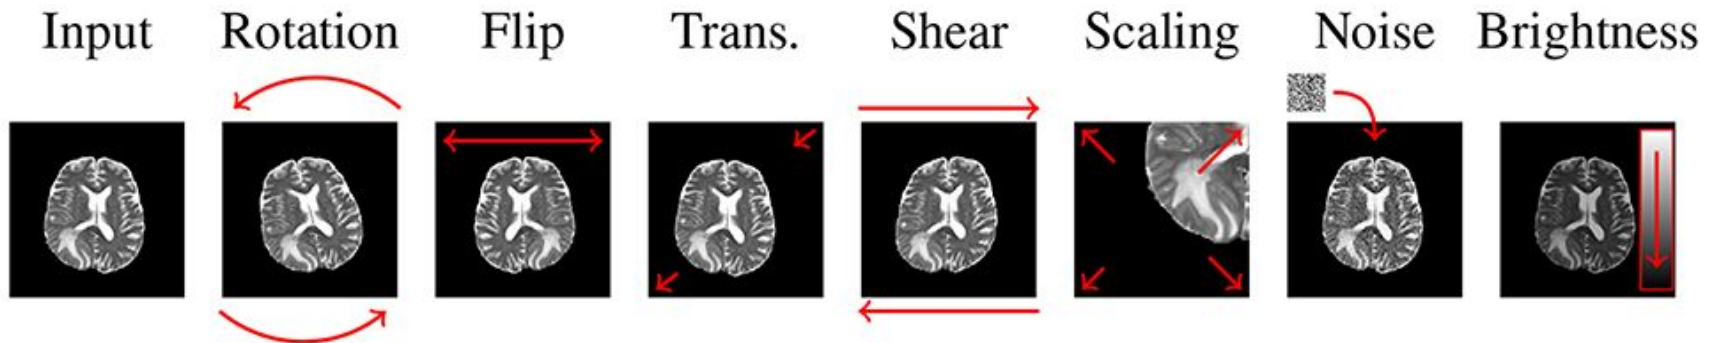
increased
training error

- more training data
- less model parameters
- L1 regularization
- L2 regularization
- dropout
- early stopping
- data augmentation

How to Train the Neural Network



We aim to **artificially increase the size of the training set** by modifying training data in a meaningful manner. Some prominent data augmentation techniques in medical image analysis:



E1: Time for Some Practical Exercises!

one epoch represents one run through all samples present in the train set



```
for epoch in range(2): # loop over the dataset multiple times
```

```
    running_loss = 0.0
```

```
    for i, data in enumerate(trainloader, 0):
```

```
        # get the inputs
```

x_i, y_i → `inputs, labels = data`

```
        # zero the parameter gradients
```

```
        optimizer.zero_grad()
```

```
        # forward + backward + optimize
```

\hat{y}_i → `outputs = net(inputs)`

```
        loss = criterion(outputs, labels) ←  $\mathcal{L}(\hat{y}_i, y_i)$ 
```

```
        loss.backward()
```

```
        optimizer.step() ←  $\theta_{k+1} = \theta_k - \alpha \nabla_{\theta} L(\theta_k, \mathbf{x}_{\{1..m\}}, \mathbf{y}_{\{1..m\}})$ 
```

```
        # print statistics
```

```
        running_loss += loss.item()
```

```
    if i % 2000 == 1999: # print every 2000 mini-batches
```

```
        print('%5d, %5d] loss: %.3f' %
```

```
              (epoch + 1, i + 1, running_loss / 2000))
```


```
        running_loss = 0.0
```

```
print('Finished Training')
```



Please have a look at your notebook and work on:

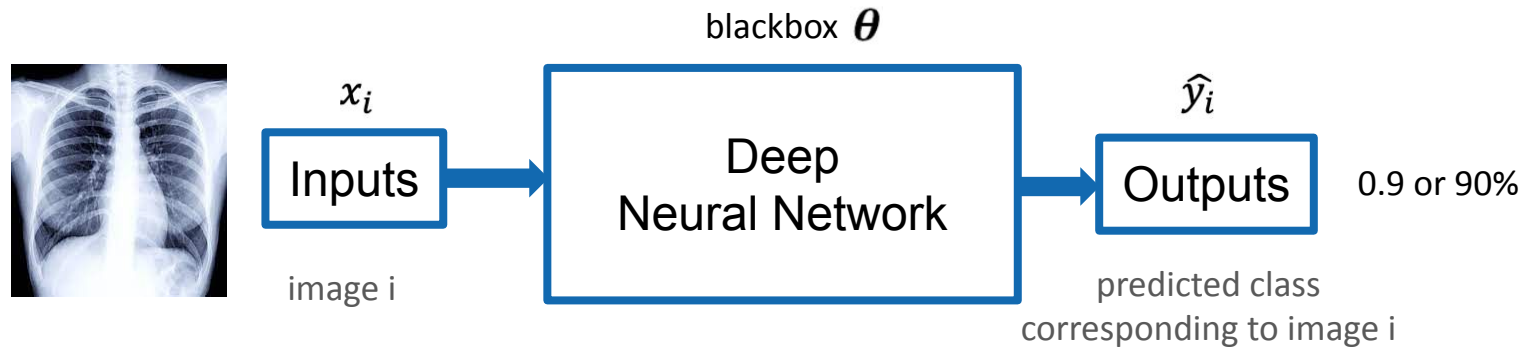
- 1) Introduction to PyTorch
- 2) Preparing Data
- 3) Implementation of a FCN/MLP



The image to the left may help you to connect the dots between the learned theory and the practical exercises you will have to tackle in the notebook.

Again, we would like to remind you to please ask whenever you have any question!


Neural Networks



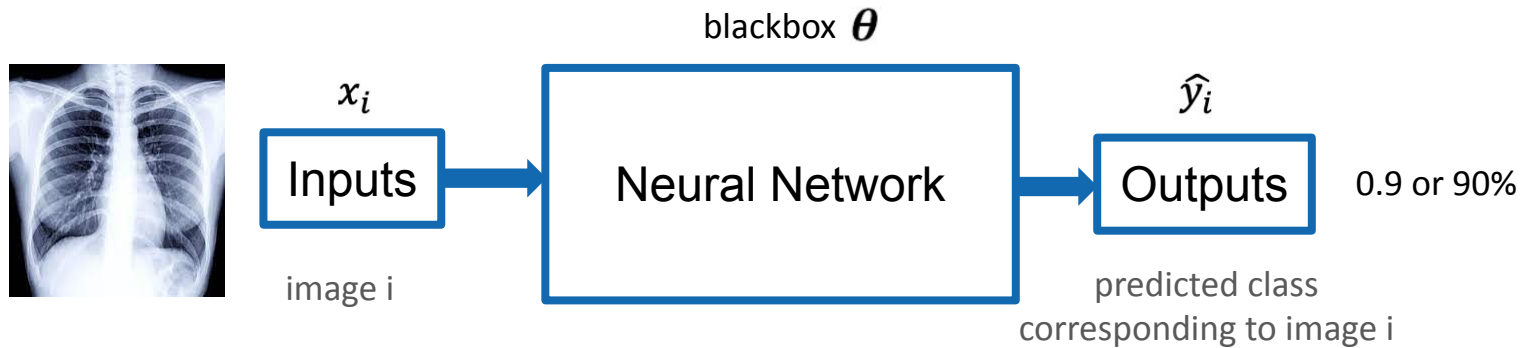
$$\hat{y}_i = W_5 \sigma(W_4 \tanh(W_3, \max(0, W_2 \max(0, W_1 x_i))))$$

lernable
parameters θ

we organize a neural network
into **multiple layers** separated by
non-linearities.




Neural Networks




$$\hat{y}_i = W_5 \sigma(W_4 \tanh(W_3, \max(0, W_2 \max(0, W_1 x_i))))$$

we organize a neural network
into **multiple layers** separated by
non-linearities.


let us explore more
advanced Neural Networks!



Neural Networks



convolution: $f * g = \int_{-\infty}^{\infty} f(\tau)g(t - \tau)d\tau$



input image

kernel weights

$$* \begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix} =$$

kernel for edge
detectionprocessed output
image

Neural Networks

 x_i

Inputs

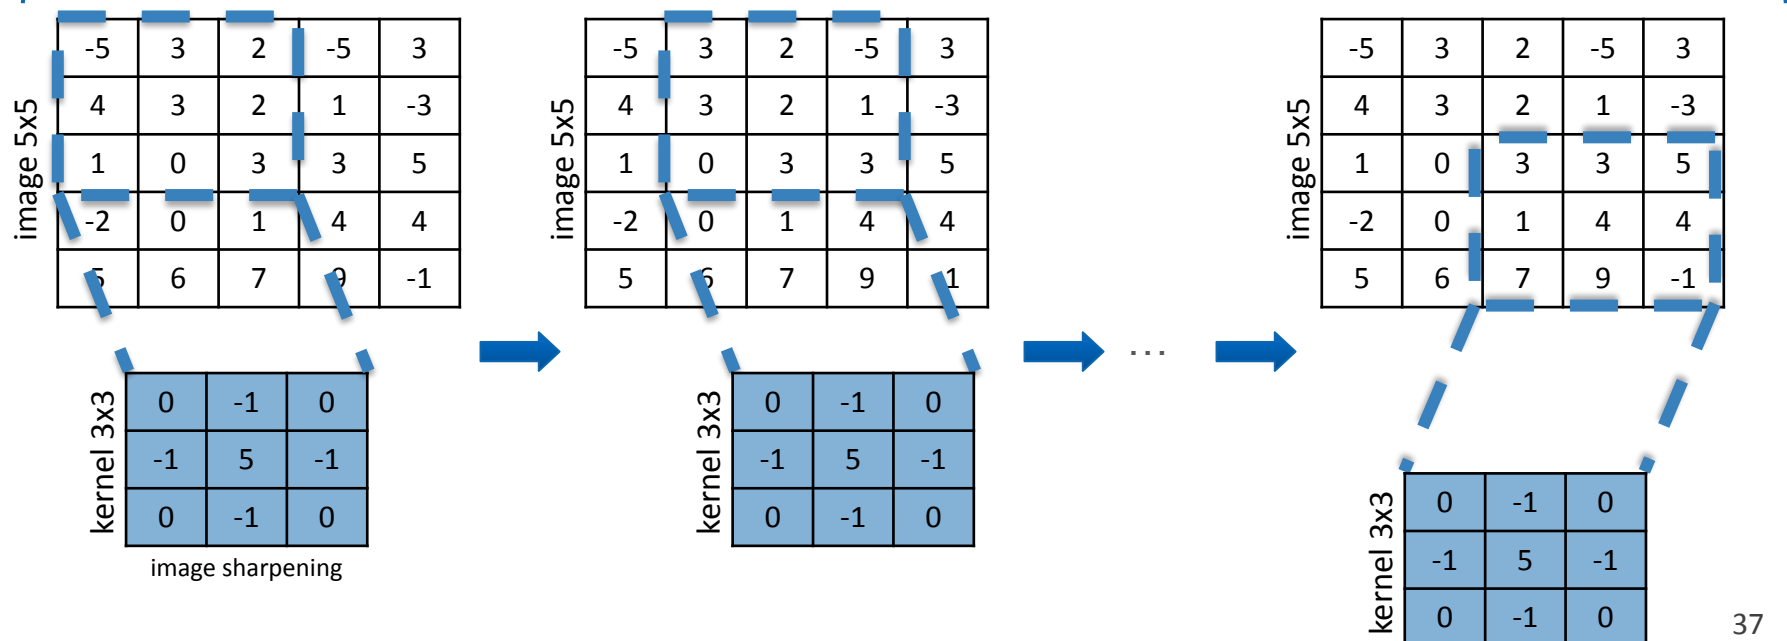
image i blackbox θ

Neural Network

 \hat{y}_i

Outputs

0.9 or 90%

predicted class
corresponding to image i 

Neural Networks

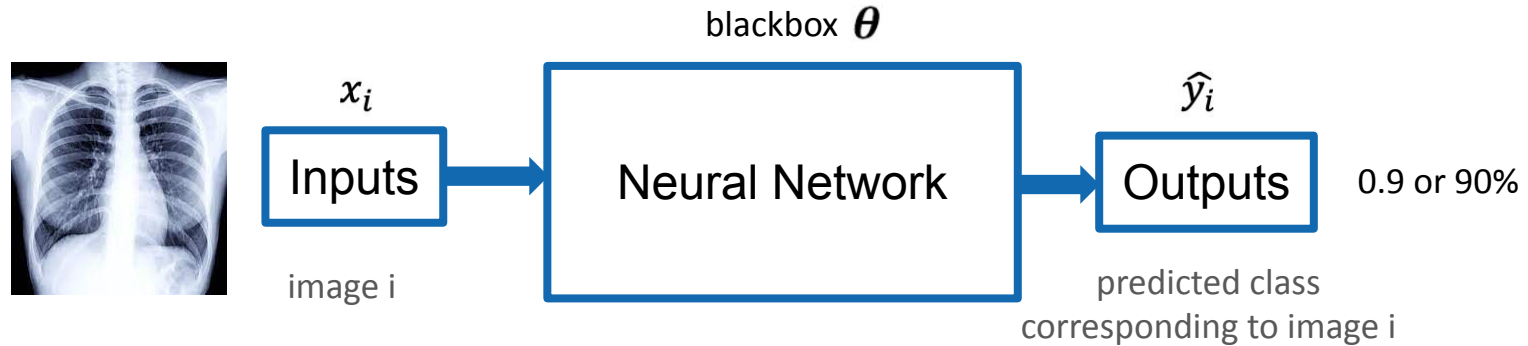


image 5x5

-5	3	2	-5	3
4	3	2	1	-3
1	0	3	3	5
-2	0	1	4	4
5	6	7	9	-1

kernel 3x3

0	-1	0
-1	5	-1
0	-1	0

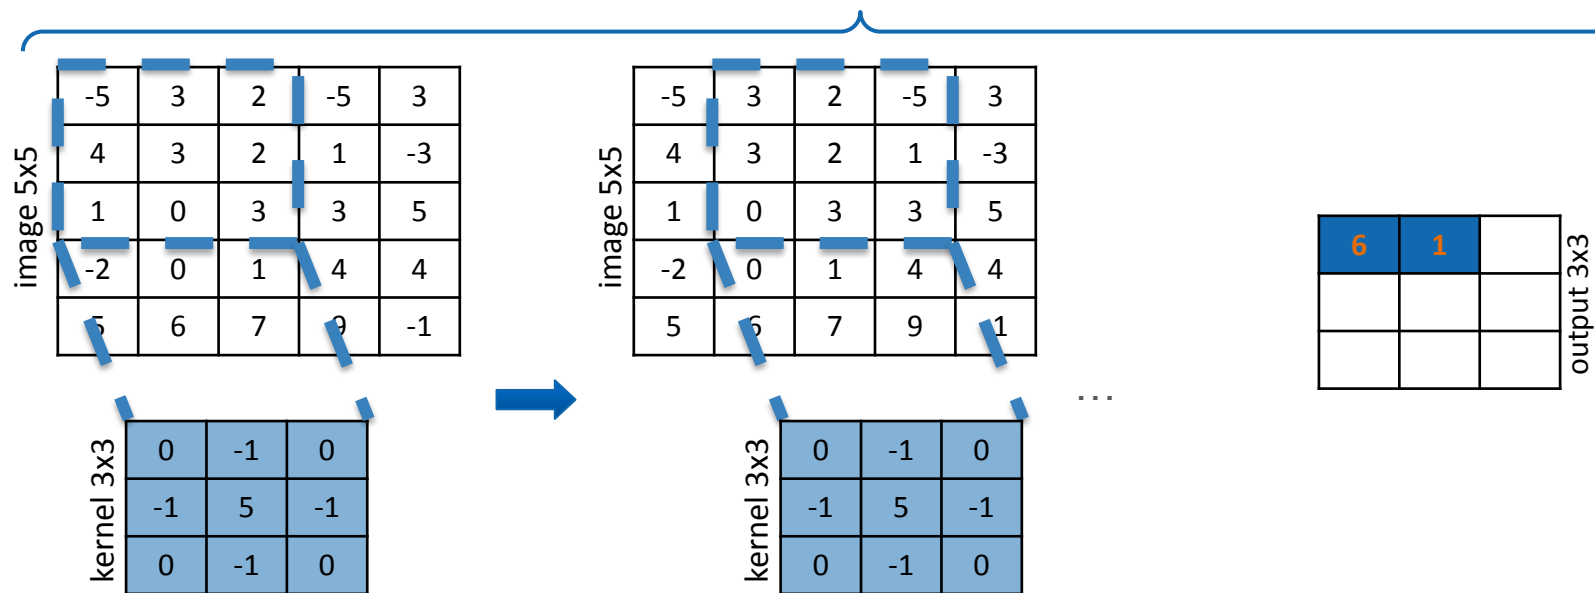
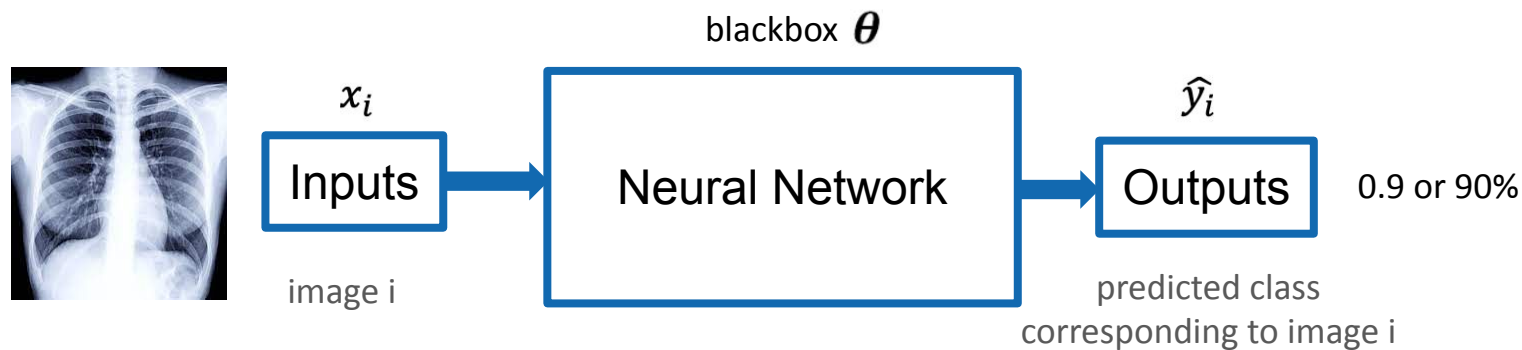
6		

output 3x3

$$5 \cdot 3 + -1 \cdot 3 + -1 \cdot 2 + -1 \cdot 0 + -1 \cdot 4$$

$$= 15 - 9 = 6$$

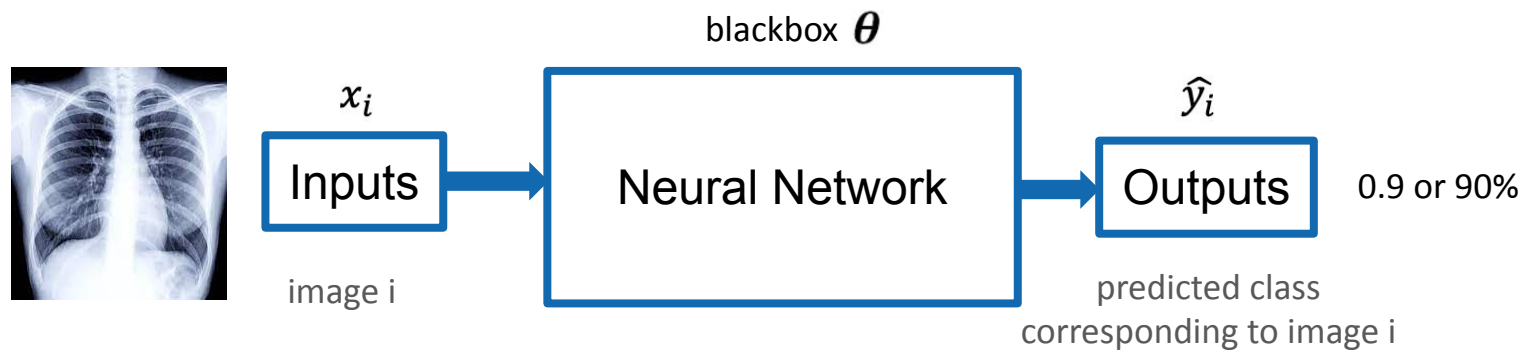
Neural Networks




$$5 \cdot 3 + -1 \cdot 3 + -1 \cdot 2 + -1 \cdot 0 + -1 \cdot 4 = 15 - 9 = 6$$

$$5 \cdot 2 + -1 \cdot 2 + -1 \cdot 1 + -1 \cdot 3 + -1 \cdot 3 = 10 - 9 - 1 = 0$$

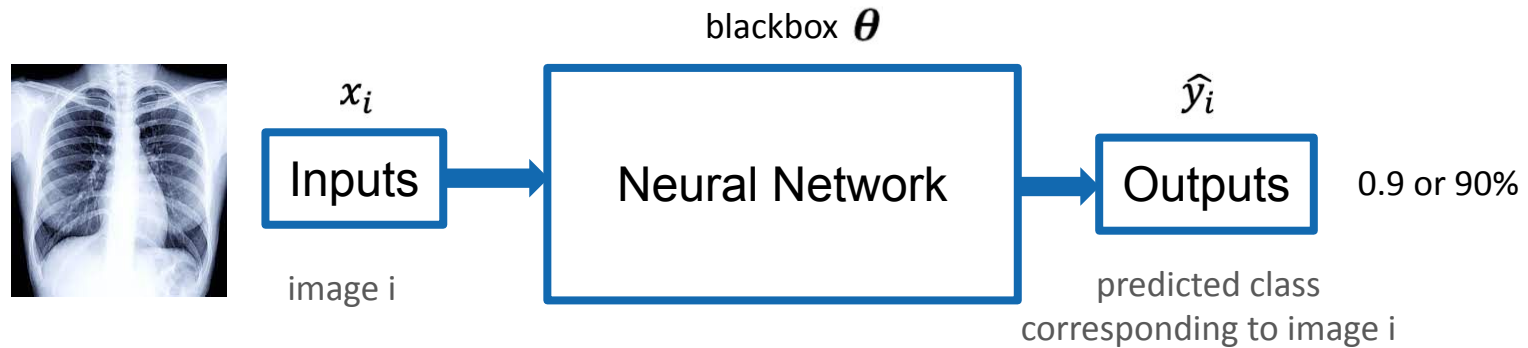
Neural Networks



filter must extend to all image channels!




Neural Networks




how is can we use this for our neural networks?

we **learn the kernel weights** and optimize them as parameters during training!

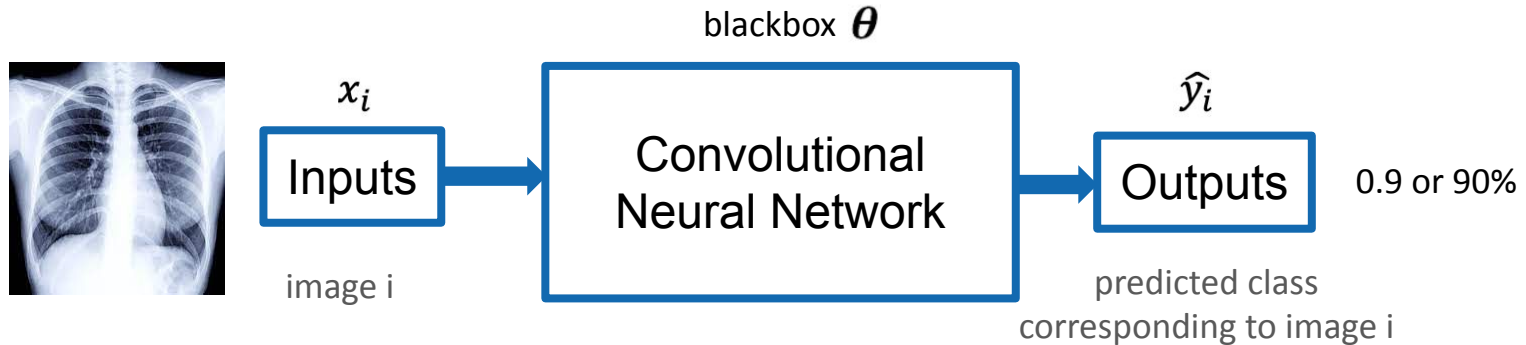


Logistic Regression:
a 1 layer, 1 neuron neural network





RGB image 3x32x32

Convolutional Neural Networks (CNN)



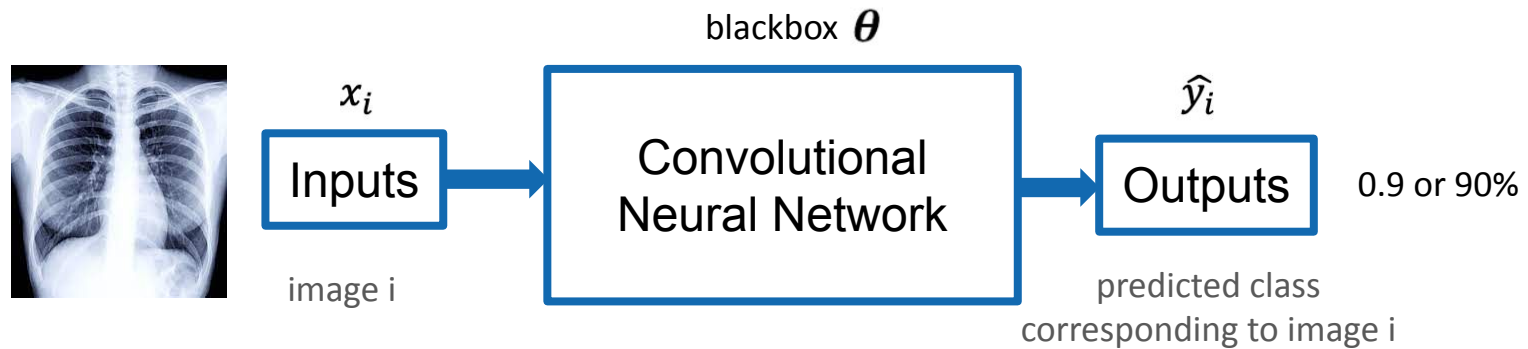
downsampling via max-pooling
(lower comp. cost; higher receptive field)



final linear layers
(for classification;
here they want to predict 1000 classes)


convolution+ReLU
 max pooling
 fully connected+ReLU

Convolutional Neural Networks (CNN)




the intuition behind CNNs:

conv1




Layer 1

conv2




Layer 2

conv3




Layer 3

conv4



Layer 4

conv5

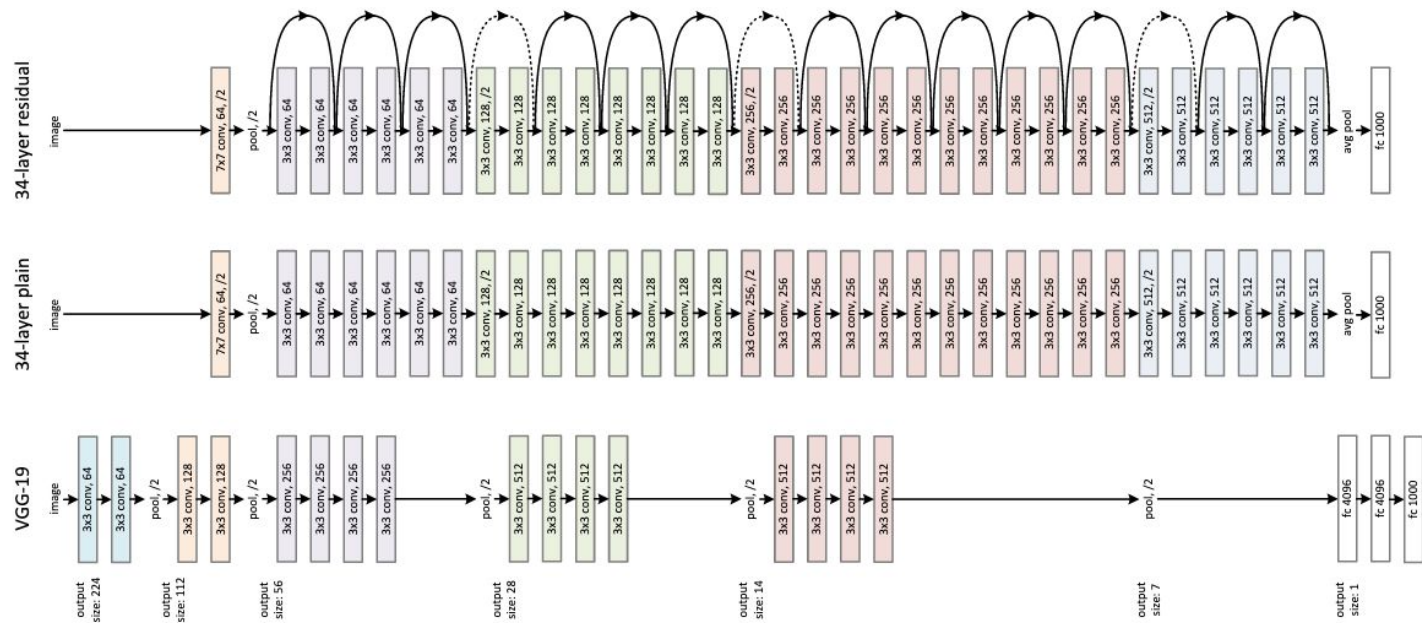
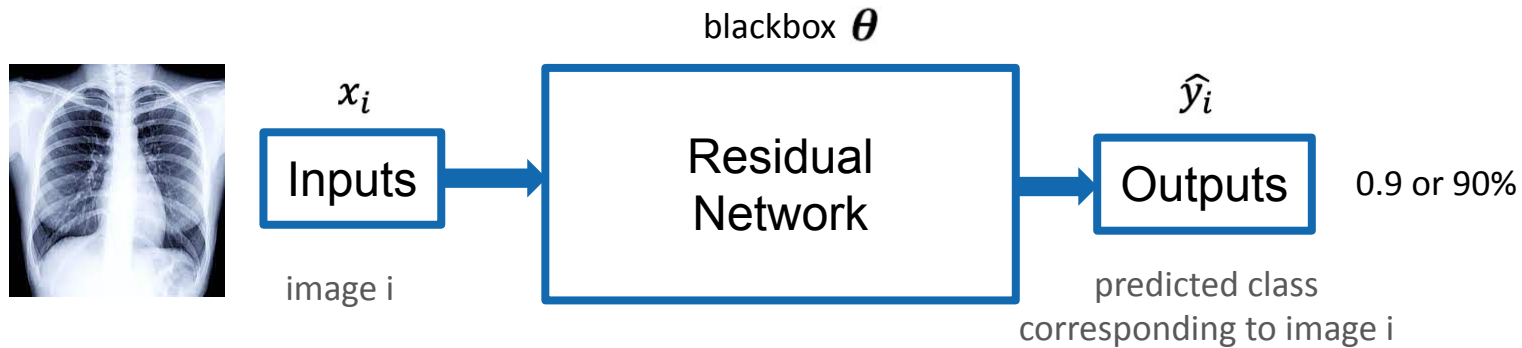


Layer 5

Zeiler, Matthew D., and Rob Fergus. "Visualizing and understanding convolutional networks." *Computer Vision—ECCV 2014: 13th European Conference, Zurich, Switzerland, September 6–12, 2014, Proceedings, Part I 13*. Springer International Publishing, 2014.

later layers learn more complex features

Residual Networks (ResNet)



a quite
deep
CNN

E2: Time for More Practical Exercises!

```
for epoch in range(2): # loop over the dataset multiple times

    running_loss = 0.0
    for i, data in enumerate(trainloader, 0):
        # get the inputs
        inputs, labels = data

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
        outputs = net(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        # print statistics
        running_loss += loss.item()
        if i % 2000 == 1999:    # print every 2000 mini-batches
            print('[%d, %5d] loss: %.3f' %
                  (epoch + 1, i + 1, running_loss / 2000))
            running_loss = 0.0

    print('Finished Training')
```



Please have a look at your notebook and work on:

- 1) Implementation of simple CNN
- 2) Experiments with ResNet
- 3) Interpretability of Predictions

Again, we would like to remind you to please ask whenever you have any question!

Please note that the last exercise on interpretability is optional and ment for students with prior experience.